




Deep Learning

Lakshmi Narasimhan
IIT Madras

Outline

1. Recap of Machine learning fundamentals

- Linear algebra and probability
- Supervised learning: Classification and Regression
- Loss and errors
- Model selection: over/under fitting, Bias-Variance tradeoff

Outline

2. Multilayer perceptron

- Function approximation theorems
- Feedforward neural networks
- Activation functions
- Width, depth and expressivity

Outline

3. Training neural networks

- Backpropagation
- Computational graph
- Optimization: Stochastic gradient descent and variants
- Avoiding overfitting: regularization, dropout, data augmentation
- Gradient instability, skip connection, initializations, early stop
- Normalization: input, layer, batch
- Hyperparameter tuning, curriculum training, transfer learning

Outline

4. Convolutional neural networks

- Multi-channel inputs
- Convolution and pooling
- Padding, striding, flattening
- Example architectures: AlexNet, GoogleNet, VGGNet, ResNet

Outline

5. Recurrent neural networks

- Recursive neural networks
- Backpropagation through time
- Neural nets with memory: LSTM and GRU
- Sequence learning and bidirectional LSTM
- Optimizing RNN: Teacher forcing, leaky units, skipping through time

Outline

6. Attention mechanism

- Self attention and multi-head attention
- Query-key-value
- Attention pooling and scoring
- Embedding and encoding
- Transformer architecture

Outline

7. Autoencoders



- Bottleneck architecture for encoder and decoders
- PCA/sparse decomposition with autoencoders
- Contractive and denoising autoencoders
- Variational autoencoder

Outline

8. State-of-the-art

- Graph neural networks
- Copresheaf topological neural networks
- Physics informed neural networks
- Kolmogorov Arnold neural networks
- Diffusion and generative models

Reference materials



Logistics

- Lectures: 5.30 PM to 7 PM on Mon & Thu (Jan 5th to Apr 10th)
- Tutorials: On alternate weeks
- Assignments: One per two weeks
 - Theory: 50%
 - Programming: 50%
- Head TAs
 - Manoj Kumar (da24s018@smail.iitm.ac.in)
 - Yuvaram Singh (da24s015@smail.iitm.ac.in)

Grading

- 25% - Mid term, 22 Feb
- 25% - End term, 25 Apr
- 20% - Assignments (4 best out of 6)
- 30% - Project
 - Finalize topic by Feb 28
 - Deadline: April 22
 - Submit: Report + codes

Sample Project Ideas

- Sentence Sentiment Classification

- <https://www.kaggle.com/datasets/jp797498e/twitter-entity-sentiment-analysis>
- <https://www.kaggle.com/datasets/kazanova/sentiment140>

- Image Super Resolution

- <https://www.kaggle.com/datasets/adityachandrasekhar/image-super-resolution>
- <https://www.kaggle.com/datasets/jesucristo/super-resolution-benchmarks>

- GenAI with Variational Autoencoders

- <https://www.kaggle.com/code/basu369victor/generate-music-with-variational-autoencoder>
- <https://www.kaggle.com/code/averkij/variational-autoencoder-and-faces-generation>

- Product Recommender System

- <https://www.kaggle.com/datasets/shivamb/fashion-clothing-products-catalog>
- <https://www.kaggle.com/datasets/lokeshparab/amazon-products-dataset>

Project Report Format

- Introduction (1 mark)
- Background and state-of-the-art (1 mark)
- Problem Statement (4 marks)
- Dataset Description (2 marks)
- Solution/Methodology/Neural Net Architecture (4 marks)
- Theoretical Analysis (2 marks)
- Data Preprocessing and Training Methodology (3 marks)
- Experimental Results (5 marks for plots & description + 5 marks for attached code)
- Inferences/Insights Obtained/Difficulties faced (3 marks)
- Template - <https://www.overleaf.com/read/cbbbwmpqctbv#397b3a>

Linear algebra and probability

Vectorspaces


- A set $V : \forall \mathbf{x}, \mathbf{y} \in V, a\mathbf{x} + b\mathbf{y} \in V$, where a and b are any scalars
 - Closed under addition and scalar multiplication
- **Subspace**: a subset of V that is closed under addition and scalar multiplication
- Any element of this vectorspace is a **vector**
 - For finite-dimensional vectorspaces (e.g., $\mathbb{R}^n$): $n \times 1$ array of numbers
 - In our study, most data samples will be represented by a vector

Orthogonality & Norms

- An inner product on $\mathbb{R}^n$: $\mathbf{x}^T \mathbf{y}$
 - Measure of component of $\mathbf{x}$ on $\mathbf{y}$ (and vice versa)
 - **Orthogonal**: $\mathbf{x}^T \mathbf{y} = 0$ (perpendicular vectors)
 - $\mathbf{x}^T \mathbf{y} > 0$: $\mathbf{x}$ and $\mathbf{y}$ have components along the same direction
 - $\mathbf{x}^T \mathbf{y} < 0$: $\mathbf{x}$ and $\mathbf{y}$ have components along opposite directions


- **p - norm** ($p \geq 1$) : $\|\mathbf{x}\|_p := \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}$

- **Euclidean norm**: $\|\mathbf{x}\|^2 = \mathbf{x}^T \mathbf{x}$



Hyperplanes

- **Linear hyperplane**
 - An $n - 1$ dimensional subspace of $\mathbb{R}^n$
 - Passes through the origin
 - Divides $\mathbb{R}^n$ into two halves
 - Represented by a normal vector
- **Normal vector:**
 - A vector orthogonal to all vectors on the hyperplane
 - Distance of $\mathbf{x}$ from hyperplane = $\mathbf{x}^T \mathbf{v} / \|\mathbf{v}\|$



Singular Value Decomposition

- Singular value decomposition : $\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$
 - Exists for any matrix
 - $\mathbf{U}, \mathbf{V}$ are orthogonal: left and right **singular vectors**
 - $\mathbf{\Sigma}$ is PSD diagonal: The diagonal elements are **singular values**
 - The diagonal elements of $\mathbf{\Sigma}$ are in ascending order : $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_n \geq 0$
- **Pseudoinverse**: $\mathbf{V}\mathbf{\Sigma}_r^{-1}\mathbf{U}^T$, where r is the rank of $\mathbf{M}$
- **Matrix norm**: $\|A\|_p = \sup_{x \neq 0} \frac{\|Ax\|_p}{\|x\|_p}$ $\|A\|_2 = \sigma_{\max}(A)$

Random Variable

- Almost all data collected in real life are affected by noise
 - This measurement is an outcome of the data collection experiment
 - The outcome may change every time the experiment is performed
 - However, certain *statistical* properties of these outcomes are deterministic
- **Probability space:** (Ω, F, P) – quantifying randomness of measurements
 - Ω : Sample space; F : Event space – all possible subsets of Ω
 - P : a measure for all elements of F – e.g., proportion of occurrence of any $u \in F$
- Random variable: a mapping from Ω to $\mathbb{R}$ (real numbers)

PMF vs PDF

Discrete random variables	Continuous random variables
Ω is countable all $w \in \Omega$ has $0 \leq P(w) \leq 1$	Ω is uncountable $u \in F$ has $0 \leq P(u) \leq 1$, $w \in \Omega$ has $P(w) = 0$
$X(w)$ is a real-discrete-valued r.v. $P(w)$ is its probability <i>mass</i>	$X(w)$ is a real-continuous-valued r.v. $P(w)$ is its probability <i>density</i>
Probability mass function (PMF) $P(X = x) = P(w)$ for $X(w) = x$	Probability density function (PDF) $P(X \in [a,b]) = \int_a^b P(x)dx = P(u)$, $u = \{X^{-1}(x), x \in [a,b]\}$
Law of total probability $\sum_i P(X = x_i) = P(\Omega) = 1$	Law of total probability $\int P(x)dx = P(\Omega) = 1$

Conditional Probability & Independence

- $P(X | Y)$: probability of occurrence of X when Y has occurred
 - The **sample space is reduced** by conditioning
 - Creates a new probability space and a valid PMF/PDF – $P(X = x | Y)$
 - Probability can increase or decrease upon conditioning
 - $P(X | Y) = P(X, Y) / P(Y)$
 - $P(X | Y = y)$ is a function of x , but $P(X | Y)$ is a function of both x and y
- **Independent** r.v.s: X and Y do not influence each other
 - $P(X, Y) = P(X) P(Y) \implies P(X | Y) = P(X)$

Bayes Theorem

- $P(X | Y) = P(Y | X) P(X) / P(Y)$
 - Given that Y is observed and a dependent hidden variable X exists, we can infer the probability of the hidden variable
- $P(X)$: Prior, statistics known before any observation
- $P(X | Y)$: Posterior, computed indirectly after observing Y
 - $P(X | Y = y)$ is a valid PMF/PDF; $P(X | Y = y) \propto P(Y = y | X) P(X)$
- $P(Y | X)$: Likelihood, computed based on the outcome of Y
 - $P(Y = y | X)$ is not a valid PMF/PDF
 - Note that the probability space itself changes for each value of X

Gaussian Distribution

- Gaussian $X \sim \mathcal{N}(\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{\sigma^2}}$
 - μ, σ : Only two parameters describe the entire PDF
- **Multivariate Gaussian:** $\mathbf{x} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^n \det \boldsymbol{\Sigma}}} e^{-(\mathbf{x}-\boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x}-\boldsymbol{\mu})}$
 - $\mathbf{x} = [X_1, \dots, X_n]^T$ has n Gaussian r.v.'s that are also **jointly** Gaussian
 - $\boldsymbol{\mu} = \mathbb{E}\{\mathbf{x}\}$ is the mean vector, $\boldsymbol{\Sigma} = \mathbb{E}\{(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T\}$ is a PSD covariance matrix
 - If $\mathbf{z} = \boldsymbol{\Sigma}^{-1/2} \mathbf{x}$, then the elements of $\mathbf{z}$ are independent Gaussian r.v.
 - Precision parameter: $\boldsymbol{\beta} = \boldsymbol{\Sigma}^{-1}$

KL Divergence

- **Maximum entropy distributions** (priors with least bias)
 - Uniform for discrete r.v.
 - Gaussian for continuous r.v.
- **Relative entropy**: $D(P \parallel Q) = \mathbb{E}_P\{\log P/Q\} \geq 0$
 - Measure of how distant is the distribution Q is different from P
- **Cross entropy**: $H(P; Q) = \mathbb{E}_P\{-\log Q\} \geq 0$ (for discrete r.v.)
 - $H(P; Q) = D(P \parallel Q) + H(P)$
 - Minimizing CE is equivalent to minimizing KLD


Classical vs Bayesian

Bayesian	Classical
Parameter to be estimated is a random variable	Parameter to be estimated is an unknown deterministic quantity
A prior distribution on the parameters is used	No prior (all values are equally likely - special case of Bayesian)
Posterior probability is optimized	Likelihood is optimized
Poor performance when prior is incorrect	Poor performance whenever the parameter is not uniformly distributed

Taylor Expansion


- Polynomial approximation

$$T(\mathbf{x}) = f(\mathbf{a}) + (\mathbf{x} - \mathbf{a})^\top Df(\mathbf{a}) + \frac{1}{2!}(\mathbf{x} - \mathbf{a})^\top \{D^2 f(\mathbf{a})\} (\mathbf{x} - \mathbf{a}) + \dots$$



- ($\nabla f(p) = \begin{bmatrix} \frac{\partial f}{\partial x_1}(p) \\ \vdots \\ \frac{\partial f}{\partial x_n}(p) \end{bmatrix}$) Hessian


$$\mathbf{H}_f = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \dots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$



Machine Learning Fundamentals

Supervised/Unsupervised Learning

Supervised Learning	Unsupervised Learning
Training + Testing (inference) phase	Only inference phase
Requires labeled data for training	All data are unlabeled
Examples: Classification & estimation	Clustering, PCA, pattern finding



Loss Functions

- How to find the best parameters while training a supervised learning model?
 - Minimize the error between model output for input x_i and the label of x_i
- Loss function: A quantification of the error obtained during parameter optimization $\theta = \operatorname{argmin} \operatorname{Loss}\{f(x, \theta) ; \operatorname{label}(x)\}$
 - **L0** norm (sparsity-promoting), 0-1 loss: $\|y - y_{\text{true}}\|_0$
 - **L1** norm, absolute error: $\|y - y_{\text{true}}\|_1$
 - **Squared L2** norm, Euclidean distance: $\|y - y_{\text{true}}\|_2^2$
 - Mixed norms: linear combination of different p -norms ($p \geq 1$) – **convex losses**
 - Cross entropy, F-score, etc.

Regression

- Setup: $y_i = f(x_i) + \varepsilon$
- Find: $f(\cdot)$. Parameterized $f(x) = g(\mathbf{w}, x)$; linear regression: $g(\mathbf{w}, \mathbf{x}) = \mathbf{x}^T \mathbf{W}$
- Given (training data)
 - Labels: $\{y_i\}$
 - Inputs : $\{x_i\}$
- Predict (testing data)
 - Compute $\{y\}$ for unseen $\{x\}$

Classification

- Setup: $x_{i,j} \in C_i$; $i = 1, \dots, K$ and $j = 1, \dots, N$
- Find: $f(.)$ s.t. $f(x_{i,j}) = C_i$
- Given (training data)
 - Labels: $\{y_i\}$ s.t. $y_i \in \{C_i\}$
 - Inputs : $\{x_{i,j}\}$
- Predict (testing data)
 - Compute $\{y\}$ for unseen $\{x\}$ s.t. $y \in \{C_i\}$

Classification

- Binary classification: $K = 2$, $C_1 = +1$ and $C_2 = -1$
- Posterior probability for binary classification is

$$\begin{aligned}\Pr(y_i = +1|x_i) &= \frac{\Pr(x_i|y_i = +1) \Pr(y_i = +1)}{\Pr(x_i|y_i = +1) \Pr(y_i = +1) + \Pr(x_i|y_i = -1) \Pr(y_i = -1)} \\ &= \frac{1}{1 + e^{-a}} = \sigma(a); \quad a = \ln \frac{\Pr(x_i|y_i = +1) \Pr(y_i = +1)}{\Pr(x_i|y_i = -1) \Pr(y_i = -1)}\end{aligned}$$

$a = \log$ posterior ratios

- The **sigmoid** function naturally arises here; converts logits to probabilities

Classification

- Posterior probability for K-class classification is

$$\begin{aligned}\Pr(y_i = C_k | x_i) &= \frac{\Pr(x_i | y_i = C_k) \Pr(y_i = C_k)}{\sum_j \Pr(x_i | y_i = C_j) \Pr(y_i = C_j)} \\ &= \frac{e^{a_k}}{\sum_j e^{a_j}} = \sigma_k(\mathbf{a}); \quad a_k = \ln \Pr(x_i | y_i = C_k) \Pr(y_i = C_k)\end{aligned}$$

- The **softmax** function naturally arises here
- Converts logits (distance from hyperplanes) to probabilities

Inference on Gaussian Data

- Binary classification

$$\Pr(\mathbf{x}_i | y_i = 1) = \frac{1}{(2\pi)^{M/2} \det(\Sigma)^{1/2}} e^{-\frac{(\mathbf{x}_i - \mu_k)^T \Sigma^{-1} (\mathbf{x}_i - \mu_k)}{2}}$$

$$\Pr(y_i = +1 | \mathbf{x}_i) = \sigma(\mathbf{x}_i^T \mathbf{w} + w_0); \quad \mathbf{w} = \Sigma^{-1}(\mu_1 - \mu_{-1})$$

$$w_0 = -\frac{\mu_1^T \Sigma^{-1} \mu_1}{2} + \frac{\mu_{-1}^T \Sigma^{-1} \mu_{-1}}{2} + \ln \frac{\Pr(y_i = 1)}{\Pr(y_i = -1)}$$

- Multiclass
classification

$$\Pr(y_i = +1 | \mathbf{x}_i) = \sigma_k(\mathbf{x}_i^T \mathbf{W} + \mathbf{w}_0); \quad \mathbf{W} = \Sigma^{-1}[\mu_1, \dots, \mu_K]$$

$$\mathbf{w}_0 = \left[-\frac{\mu_1^T \Sigma^{-1} \mu_1}{2} + \ln \Pr(y_i = 1), \dots, -\frac{\mu_K^T \Sigma^{-1} \mu_K}{2} + \ln \Pr(y_i = K) \right]$$

- Ridge regression


$$\hat{\boldsymbol{\theta}} = \arg \min_{\boldsymbol{\theta}} \|\mathbf{y} - \mathbf{G}\boldsymbol{\theta}\|_2^2 \quad \text{s.t.} \quad \|\boldsymbol{\theta}\|_2^2 \leq B$$

$$= \arg \min_{\boldsymbol{\theta}} \|\mathbf{y} - \mathbf{G}\boldsymbol{\theta}\|_2^2 + \lambda \|\boldsymbol{\theta}\|_2^2$$

$$= (\lambda I + \mathbf{G}^T \mathbf{G})^{-1} \mathbf{G}^T \mathbf{y}$$

Classifying Linearly Separable Data

- Hyperplane projection: $(\mathbf{x} - \mathbf{w}_0)^T \mathbf{w} = \mathbf{x}^T \mathbf{w} - \mathbf{w}_0^T \mathbf{w} = \mathbf{x}^T \mathbf{w} - w_0$
 - $\mathbf{w}_0$ is the **shift of origin** and $\mathbf{w}$ is **normal to the hyperplane**
- Binary classification
 - When $\mathbf{x}$ is on the hyperplane, $\Pr(y = C_1) = \sigma(0) = 1/2$
 - Farther from the hyperplane $\sigma(+ve) \rightarrow 1$ or $\sigma(-ve) \rightarrow 0$
 - Sigmoid output gives the PMF of the **Bernoulli** r.v. $f(x)$
- Multi-class classification
 - When $\mathbf{x}$ is on the hyperplane, $\Pr(y = C_1) = \sigma_k(\mathbf{0}) = 1/K$
 - Farther from the k^{th} hyperplane $\sigma_k(+ve) \rightarrow 1$ or $\sigma_k(-ve) \rightarrow 0$
 - Softmax output gives the PMF of the **categorical** r.v. $f(x)$



Logistic Regression – Single Layer Neural Net

- Let the classes be **one-hot encoded** (to represent a PMF)

$$\Pr(\mathbf{y}|\mathbf{x}, \mathbf{W}) = \prod_{k=1}^K \Pr(y = C_k|\mathbf{x}, \mathbf{W})^{y_k}; \quad \sum_{k=1}^K y_k = 1; \quad \sum_{k=1}^K \Pr(y = C_k|\mathbf{x}, \mathbf{W}) = 1$$

- To find $\mathbf{w}$: do **maximum likelihood regression**

$$\arg \min_{\mathbf{W}} -\ln \Pr(\mathbf{y}|\mathbf{x}, \mathbf{W}) = \arg \min_{\mathbf{W}} - \sum_{i=1}^N \sum_{k=1}^K y_{i,k} \ln \sigma_k(\mathbf{x}_i^T \mathbf{W}) = \arg \min_{\mathbf{W}} \sum_i H(y_i; \sigma_k(\mathbf{x}_i^T \mathbf{W}))$$

- The optimal hyperplanes are obtained using gradient descent




$$\mathbf{w}_{j,t+1} = \mathbf{w}_{j,t} - \eta \nabla_{\mathbf{w}_j} \mathcal{L} = \mathbf{w}_{j,t} - \eta \sum_i (\sigma_j(\mathbf{x}_i^T \mathbf{W}_t) - y_{i,j}) \mathbf{x}_i$$

Summary

- We focus on parametric supervised learning framework
- Neural models can be viewed as interpolators or curve-fitters
- Loss function: measure of deviation from expected output/label
- Optimal linear/ridge regression: a single layer linear neural net
- Optimal linear classification: a single layer neural net with sigmoid/softmax activation
- Learning weights from training data implicitly learns data prior


Extending to Multiple Layers




- We cannot model XOR with a linear classifier
- What if cascade multiple single layer neural net?
 - Similar to making intermediate decisions
- Can we keep increasing the number of layers/parameters?



Training/Generalization Error

- Increasing number of parameters is not always good
 - Training error may go down, but testing (generalization) error may be high
- There is an optimal number of parameters for a given kind of data
 - VC dimension ($|\text{Test error} - \text{Training error}| \leq O(d_{VC} \log n/n)^{1/2}$)
- Underfitting: Less than optimal
- Overfitting: More than optimal
 - How to identify overfitting?






Validation Error & Overfitting

- Split available labeled data to training + validation sets
 - Sequentially minimize both training and validation errors
 - Use validation set to identify overfitting

- Validation error: The error measured with validation set – An indicator of generalization error




- **Cross validation:** When very few labeled data is available
 - Use different training+validation split in each iteration of training

Bias-Variance & Model-fitting

- Bias: $\mathbb{E}\{y - y_{\text{true}}\}$ – average deviation from true value
- Variance: $\text{Var}\{y\}$ – average deviation of inference
- Is lower bias or variance always good?

$$\begin{aligned}\mathbb{E}((y - \hat{y})^2) &= \mathbb{E}(y^2) + \mathbb{E}(\hat{y}^2) - 2\mathbb{E}(\hat{y})\mathbb{E}(y) \\ &= \mathbb{E}(y^2) - \mathbb{E}(y)^2 + \mathbb{E}(\hat{y}^2) - \mathbb{E}(\hat{y})^2 - 2\mathbb{E}(\hat{y})\mathbb{E}(y) + \mathbb{E}(y)^2 + \mathbb{E}(\hat{y})^2 \\ &= \text{Var}(y) + \text{Var}(\hat{y}) + \text{Bias}^2\end{aligned}$$



- Overfitting : High model variance and low MSE in training
- Underfitting : Low model variance and high MSE in training

Bayesian Inference


- Infer bias (p) of a coin from observations of its toss: $X_i \sim \text{Be}(p)$
- (Classical) No bias

$$\hat{p}_1 = \arg \max_p p^{\sum y_i} (1 - p)^{N - \sum y_i} = \frac{\sum y_i}{N}$$

$$\text{Var}(\hat{p}_1) = \frac{p(1 - p)}{N} \quad \text{Bias}(\hat{p}_1) = 0$$

$$\hat{p}_3 = \arg \max_p p^{\sum y_i} (1 - p)^{N - \sum y_i} P(p) = \frac{\sum y_i}{N + 1}$$

- (Bayesian) Bias is already in-built in prior




No Free-Lunch Theorem

- Training dataset may not contain all possible samples of input
 - How would the unseen data affect the learned statistics?
- **No free-lunch theorem**: As the unseen data can have different statistics, there can exist a dataset where any model (optimally trained with a training dataset) can perform poorly
 - So, will all models perform equally bad on average?
 - Tradeoff between **robustness** and **accuracy**
- **Inductive Bias**: If some statistics of the entire data is known, use that information in addition to training data to optimize the model – **different architectures**


Double Descent

- Overfitting can be overcome
 - With deeper networks
- As the number of parameters (depth) increases, the overfitting is smoothened
 - Smoother interpolations and extrapolations
- Interpolation threshold
 - Where the number of parameters equals the number of data points
 - R. Shaeffer et al, *Double Descent Demystified: Identifying, Interpreting & Ablating the Sources of a Deep Learning Puzzle*
 - Some times more training data can hurt performance!



Grokking

- Generalization of over-parameterized nets by **over training**
- Amount of over training reduces with decrease in dataset size
 - *Grokking Explained: A Statistical Phenomenon, 2025*
- **Weight decay** helps
 - Additional term in the loss function to push the weights closer to zero



Assignment

- Generate training 100 samples : $y = x + 0.3x^2 + 0.7x^3 + w$
 - x is uniformly distributed in $[-3 \text{ to } 3]$
 - w is Gaussian noise with mean 0 and standard deviation 0.09
- Perform linear regression for model sizes $M = 1$ to 120 and plot
 - Bias (average of $y - f(x)$) vs M
 - Variance of y vs M
 - Training error vs M
 - Validation error vs M

Multilayer Perceptron

Extending Single-Layer Neural Nets

- For linear data models
 - $y = \mathbf{w}^T \mathbf{x}$ or $\sigma(\mathbf{w}^T \mathbf{x} + b)$
 - Perceptron: $u(\mathbf{w}^T \mathbf{x} + b)$
- In general, $y = f(\mathbf{w}, \mathbf{x})$
 - What are some good parametric functions?
 - Are there a family of such functions that can fit any arbitrary N -dimensional curve?

Universal Function Approximation Theorems



- Stone - Weierstrass (1885)
 - Every continuous function can be approximated by polynomials
- Andrew Barron (1993 - Universal Approximation Bounds for Superpositions)
 - Lipschitz functions can be approximated by infinite-width 1-layer NN
- Cybenko (1989 - Approximation by superpositions of a sigmoidal function)
 - Lipschitz functions can be approximated by infinite-width 1-layer sigmoid NN

Universal Function Approximation Theorems

- Hornik – Stinchcombe – White (1989 - Multilayer feedforward networks are universal approximators)
 - Lipschitz functions can be approximated by finite-width multi-layer NN
- Lu – Pu – Wang – Hu – Wang (2017 - The expressive power of neural network)
 - Lipschitz functions can be approximated by finite-width multi-layer ReLU NN
- Kolmogorov–Arnold (1957)
 - Lipschitz functions can be approximated by 2-layer NN

Universal Function Approximation


- Polynomials in high dimensions : very high complexity
- Any function $f(x)$ (with some regularity conditions) can be approximated by a superposition of scaled and shifted step functions
 - $f(x) = \sum a_i u(w_i x - b_i)$
 - w_i & b_i are the **scale** and **shift** factors
 - $u()$ is the **activation function**



Example


- Fit $w_1 u(x + w_0)$ with MSE loss

$$\text{Loss} = (0 - w_1 u(0 + w_0))^2 + (2 - w_1 u(2 + w_0))^2$$





- How does this loss function behave w.r.t. the weights?
- Is it differentiable w.r.t. the weights?

Loss Surface



Loss Surface *Smoothens* with More Data



Challenges

- How to find optimal values?
 - Find roots of derivative expression
- Derivatives should exist
- Roots: Move along the curve to find a flat region
 - Convex surface? Unique solution? Local minima?
 - Initialization point?

Activation Functions



- Step function has derivative = infinity at zero-input

- Options: Sigmoid, softmax, tanh (odd symmetry)









- Non-limiting activations:

- ReLU (rectified linear unit) & SeLU (scaled exponential linear unit)

$$\text{selu}(x) = \lambda \begin{cases} x & \text{if } x > 0 \\ \alpha e^x - \alpha & \text{if } x \leq 0 \end{cases}$$



Some Common Activation Functions

Function	Expression	Plot	Derivative
Sigmoid	$\frac{1}{1 + \exp(-x)}$		
Tanh	$\frac{1 - \exp(-2x)}{1 + \exp(-2x)}$		
ReLU	$\max(x, 0)$		
Softmax	$\frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$		


Training Neural Networks

Gradient Descent

- How to find the roots of the derivative of the loss function?
- Taylor approximation in higher dimension
 - **Gradient**: Linear approximation of a function at a point
 - **Hessian**: Curvature of a function at a point
- $L(\mathbf{w} + \boldsymbol{\varepsilon}) \approx L(\mathbf{w}) + \nabla L(\mathbf{w})^T(\mathbf{w} + \boldsymbol{\varepsilon} - \mathbf{w})$
 - For $L(\mathbf{w} + \boldsymbol{\varepsilon}) - L(\mathbf{w}) < 0$, pick $\boldsymbol{\varepsilon}$ opposite to $\nabla L(\mathbf{w})$

Gradient Descent

- Gradient descent: $\mathbf{w}_{n+1} = \mathbf{w}_n - \eta \nabla L(\mathbf{w}_n)$
 - $L(\mathbf{w}_0) \geq L(\mathbf{w}_1) \geq \dots \geq L(\mathbf{w}_n) \geq \dots$
 - η : step size or **learning rate**
- How to find gradient w.r.t to all weights?
- What is a good starting $\mathbf{w}_0$?
- What is a good learning rate?



Backpropagation

- For input-label pair (x, y) and loss function $C(\cdot)$

$$C(y, f^L(W^L f^{L-1}(W^{L-1} \dots f^2(W^2 f^1(W^1 x)) \dots)))$$

- The gradient is $\frac{dC}{da^L} \cdot \frac{da^L}{dz^L} \cdot \frac{dz^L}{da^{L-1}} \cdot \frac{da^{L-1}}{dz^{L-1}} \cdot \frac{dz^{L-1}}{da^{L-2}} \dots \frac{da^1}{dz^1} \cdot \frac{\partial z^1}{\partial x}$

$$\nabla_x C = (W^1)^T \cdot (f^1)' \circ \dots \circ (W^{L-1})^T \cdot (f^{L-1})' \circ (W^L)^T \cdot (f^L)' \circ \nabla_{a^L} C$$

- At each layer,

$$(f^1)' \circ (W^2)^T \cdot (f^2)' \circ \dots \circ (W^{L-1})^T \cdot (f^{L-1})' \circ (W^L)^T \cdot (f^L)' \circ \nabla_{a^L} C$$

$$(f^2)' \circ \dots \circ (W^{L-1})^T \cdot (f^{L-1})' \circ (W^L)^T \cdot (f^L)' \circ \nabla_{a^L} C$$

$$(f^{L-1})' \circ (W^L)^T \cdot (f^L)' \circ \nabla_{a^L} C$$

$$(f^L)' \circ \nabla_{a^L} C,$$

Backpropagation with 2-layers

$$\text{Loss} = L\{f_2(\underbrace{\mathbf{W}_2}_{=\mathbf{b}_2} \underbrace{f_1(\underbrace{\mathbf{W}_1 \mathbf{x}}_{=\mathbf{a}_1})}_{=\mathbf{a}_2}); y\}$$

$$\frac{\partial L}{\partial \mathbf{W}_2} = \mathbf{b}_1 \left(f'_2 \cdot \frac{\partial L}{\partial \mathbf{b}_2} \right)^T$$
$$\frac{\partial L}{\partial \mathbf{W}_1} = \mathbf{x} \left(f'_1 \cdot \mathbf{W}_2^T f'_2 \cdot \frac{\partial L}{\partial \mathbf{b}_2} \right)^T$$

Backpropagation

- Compute $\mathbf{g} = \nabla_{\text{output}} L$
- At layer i , (starting from the last layer)
 - $\mathbf{g} = f'_i \mathbf{W}_{i+1}^T \mathbf{g}$
 - Gradient = (current layer input) $\mathbf{g}^T$
 - Repeat till first layer
- How to get $\nabla_{\text{output}} L$?

Initialization

- Can all weights be initialized to 0?
- Why we need random initialization?
- Example: 2 output nodes and 1 input node
 - $y_1 = f(w_1x + b_1)$
 - $y_2 = f(w_2x + b_2)$
 - Gradients are similar

Initialization

- Control variance of randomness
 - Xavier: For tanh, sigmoid - $\mathcal{N}(0, 2/n_{l+1}+n_{l-1})$
 - $\text{Var}(\text{output at } l\text{th layer}) = n_{\text{input}} \sigma_w \text{Var}(\text{output at } l-1\text{th layer})$
 - $\text{Var}(\text{grad at } l\text{th layer}) = n_{\text{output}} \sigma_w \text{Var}(\text{grad at } l+1\text{th layer})$
 - Kaiming He: For ReLU - $\mathcal{N}(0, 2/n_{l-1})$
 - Almost half of the output of ReLU is zero
 - LeCun: For SELU - $\mathcal{N}(0, 1/n_{l-1})$
 - Random orthogonal matrix

AutoGrad


- How to compute gradients of arbitrary models & losses?
- Gradient of a function (in programming) can be computed using **autograd**
 - Exact gradient, not numerical gradient
 - No cycles
 - AG Baydin et al, *Automatic Differentiation in Machine Learning: a Survey*
- **Forward pass**: compute the output value of all nodes
- **Backward pass**: compute the gradients using the forward pass values and chain rule of derivatives
- Implemented using **computational graphs**

Computational Graph


- How to compute gradients of arbitrary models & losses?
- Computational graph: a **directed acyclic graph** representing all the computations between input and output
 - **Dynamic** (Loops and conditional statements)
 - **Static**
 - Output of each (intermediate) computation is a variable or node
 - Function of input variables
 - No cycles (no variable is reused in the graph)

Example Computational Graph

```
def function(x):  
    if x>0:  
        y = 1  
    else:  
        y = -1  
    z = 2*x  
    for i in range(2):  
        y = z*y + i  
    return y
```



Example Computational Graph



$$\frac{\partial f}{\partial b} = 1$$

$$\frac{\partial b}{\partial z} = a$$

$$\frac{\partial a}{\partial z} = y$$

$$\frac{\partial z}{\partial x} = 2$$



$$\frac{\partial b}{\partial a} = z$$

$$\frac{\partial a}{\partial y} = z$$

$$\frac{\partial y}{\partial x} = 0$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial b} \left[\frac{\partial b}{\partial z} \cdot \frac{\partial z}{\partial x} + \frac{\partial b}{\partial a} \cdot \left(\frac{\partial a}{\partial z} \cdot \frac{\partial z}{\partial x} + \frac{\partial a}{\partial y} \cdot \frac{\partial y}{\partial x} \right) \right] = 2(a + yz)$$

Backpropagation Example



Backpropagation

- Autograd computes gradient through backpropagation
 - Gradients *flow* from one node to another
 - Recursion and matrix operations can have derivatives
 - M. Seeger et al, Auto-Differentiating Linear Algebra, 2019
- Gradients do not flow: **vanishing** or **exploding** cases – finite precision
- Autograd fails for
 - Non-differentiable functions
 - Improper computational graph
 - e.g., `if x = 2 then 4 else x * x`, autograd output at 2 = 0 instead of 4
 - Python: array assignment (<https://github.com/HIPS/autograd/blob/master/docs/tutorial.md>)

Learning NN parameters $\mathbf{w}$ – Training Phase


- Initialize $\mathbf{w}_0$: Xavier, K. He, LeCun, etc., and compute loss -- *forward pass*
- For $n = 1$ to till convergence; Dataset with N samples,
 - Compute/estimate gradient: $\mathbf{g}_n = \sum_i \nabla_{\mathbf{w}_{n-1}} L_i / N$
 - For each i , compute $\nabla_{\mathbf{w}_{n-1}} L_i$ using backpropagation (autograd) -- *backward pass*
 - SGD, SAG, SAGA
 - Compute learning rate: η
 - AdaGrad, RMSprop, Adam
 - Update parameters: $\mathbf{w}_n = \mathbf{w}_{n-1} - \eta \mathbf{g}_n$
 - First order, second order
 - Compute loss: $L(\text{label}, \text{Model}\{\mathbf{w}_{n+1}, \text{input}\})$ -- *forward pass*
- End

Stochastic Gradient Descent

- What if number of samples is large? Compute N gradients in each step?
- SGD – at each descent step, find gradient for only one random sample:
 - $\mathbf{w}_n = \mathbf{w}_{n-1} - \eta \nabla_{\mathbf{w}} L_j$
 - j is randomly chosen from $\{1, \dots, N\}$ at every descent step:
 - Estimated gradient is a random (stochastic) sample: $\nabla_{\mathbf{w}} L_j = \nabla_{\mathbf{w}} L + z = \text{average (true) gradient} + \text{noise}$
 - Note that, for large N , $\sum_i \nabla_{\mathbf{w}} L_i / N \rightarrow \mathbb{E}\{\nabla_{\mathbf{w}} L\}$
 - For correlated data samples, $\nabla_{\mathbf{w}} L_j$ could be close to average
 - Uncorrelated samples (noisy gradient) may lead to divergence (variance of z is high)
 - Very slow convergence $O(1/\sqrt{n})$

Stochastic Gradient Descent - Batched


- Improve average gradient estimate through mini-batches
 - SGD with minibatch: $\mathbf{w}_n = \mathbf{w}_{n-1} - \eta \sum_{k=1 \text{ to } M} \nabla_{\mathbf{w}} L_{j_k} / M$
 - j_k is randomly chosen from $\{1, \dots, N\}$, for $k = 1$ to M , at every descent step
 - Batch size: $M \ll N$, reduces variance; converges at $O(1/n)$ – theoretical best
- **Epoch**: Set of descent steps required to go through all N samples
 - GD: 1 epoch = 1 descent step
 - SGD: 1 epoch = N descent steps
 - SGD-minibatch: 1 epoch = N/M descent steps




Stochastic Average Gradient

- Can we improve SGD without a mini-batch?
 - SAG: $\mathbf{w}_{n+1} = \mathbf{w}_n - \eta(\nabla_{\mathbf{w}}L_j^{(n)} - \nabla_{\mathbf{w}}L_j^{(n-1)} + \sum_i \nabla_{\mathbf{w}}L_i^{(n-1)})/N$
 - At step $n = 0$, set $\nabla_{\mathbf{w}}L_i^{(0)} = 0$ and compute average gradient as $\sum_i \nabla_{\mathbf{w}}L_i^{(0)}/N$
 - At step n , j is randomly chosen from $\{1, \dots, N\}$ and only $\nabla_{\mathbf{w}}L_j^{(n)}$ is evaluated
 - Update average gradient as $(\nabla_{\mathbf{w}}L_j^{(n)} + \sum_{i \neq j} \nabla_{\mathbf{w}}L_i^{(n-1)})/N$
 - For all i , at each step, store gradients $\nabla_{\mathbf{w}}L_i^{(n)}$
 - However, this average gradient is a biased estimate
- SAGA (SAG Accelerated): $\mathbf{w}_{n+1} = \mathbf{w}_n - \eta(\nabla_{\mathbf{w}}L_j^{(n)} - \nabla_{\mathbf{w}}L_j^{(n-1)} + \sum_i \nabla_{\mathbf{w}}L_i^{(n-1)})/N$
 - Unbiased estimate of the average gradient; converges at $O(1/n)$
 - Defazio et al "SAGA: A Fast Incremental Gradient Method With Support for Non-Strongly Convex Composite Objectives"


Effect of Learning Rate



Large step size




Small step size



Gradient

Optimal Learning rate – Steepest Descent

- Optimal $\eta^* = \operatorname{argmin}_{\eta} L(w_{n-1} - \eta \nabla w_{n-1} L) = \operatorname{argmin}_{\eta} L(w_{n-1} - \eta g_n)$
 - Not always computable
- At optimal η^* , $\nabla_{\eta} L(w_{n-1} - \eta g_n) = 0$
 - That is, $\nabla L(w_n)^T g_n = g_{n+1}^T g_n = 0$
- Consecutive gradients are orthogonal
 - Slow convergence, undoing previous minimizations
 - Fix: differently adapting learning rate or gradient



Adapting Learning Rate

- Linear decay: $\eta_n = (1 - \alpha)\eta_0 + \alpha\eta_t$
 - t: terminal rate instance; $\alpha = \min(1, n/t)$
- AdaGrad: $\eta_n = \alpha / (\epsilon + \sqrt{h_n})$; $h_n = h_{n-1} + (\sum_i \nabla_w L_i^{(n)} / N)^2$
 - Change η depending on the **magnitude of slope** of each coordinate in the gradient
 - Element-wise learning rate, ϵ is for numerical stability
 - Reduce learning rate when **accumulated gradient** is large and vice versa
 - (+) Good for sparse gradients
 - (–) decreases learning rate early – vanishing; problem when gradient is noisy
- RMSprop: $\eta_n = \alpha / (\epsilon + \sqrt{h_n})$; $h_n = \beta h_{n-1} + (1 - \beta)(\sum_i \nabla_w L_i^{(n)} / N)^2$
 - Accumulation is averaged to prevent early decrease in learning rate

Adaptive Moment Estimation



- AdaGrad & RMS prop compute root sum square & RMS of gradients, respectively
 - Results in learning rate decay as n increases and gradient magnitude accumulates
 - However, the actual direction of gradient is not used to increase or decrease learning rate
- Compute **first moment**: $\mathbf{m}_n = \beta_1 \mathbf{m}_{n-1} + (1 - \beta_1) \mathbf{g}_n$
 - β_1 is less than 1 (usually very close to 1)
 - $\mathbf{m}_n$ is **weighted moving average** of gradients – **smoother** trajectory (thus, skips small local minima)
 - For $\mathbf{m}_0 = 0$ and a constant gradient $\mathbf{g}_n = \mathbf{g}$, $\mathbf{m}_n \rightarrow \beta_1 \mathbf{g}_n$
 - Biased towards zero (as initial point is 0); **unbiased** moment = $\mathbf{m}_n / (1 - \beta_1^n)$
 - Builds **momentum** that helps to move out of small wells

Adam Optimizer


- Large momentum can lead to moving past a significant local minima
- A solution: When the variance of the gradient is high, take smaller steps to not miss local minima and vice versa
- Compute **non-central second moment**: $v_n = \beta_2 v_{n-1} + (1 - \beta_2) g_n^2$
 - β_2 is less than 1 (usually very close to 1)
 - v_n is **weighted moving average** of the (uncentered) variance of the gradients
 - Biased towards zero; **unbiased** moment $s_n = \sqrt{v_n / (1 - \beta_2^n)}$
 - Divide learning rate by $(s_n + \epsilon)$
 - Helps to **reduce oscillations**

Adam Optimizer

- Adam
 - Compute m_n and $k_n = m_n / (1 - \beta_1^n)$
 - Compute v_n and $s_n = \sqrt{v_n / (1 - \beta_2^n)}$
 - Update $\theta_{n+1} = \theta_n - \eta k_n / (s_n + \epsilon)$
 - ϵ – to avoid division by zero
 - D. P. Kingma et al, *Adam: A Method for Stochastic Optimization*
- Learning rate is adapted coordinate-wise
- Multiple variants of Adam exist in literature



Adam Method



Nesterov Optimization

- Evaluate gradient near the update point
 - Predict the next point and find the gradient there
- Compute $\nabla_{\mathbf{w}_{n-1} + \mu \mathbf{r}_{n-1}} L_i$ instead of $\nabla_{\mathbf{w}_n} L_i$
 - $\mathbf{r}_n$ is the parameter update term such that $\theta_n = \theta_{n-1} + \mathbf{r}_n$
 - Convergence rate: $O(1/n^2)$
- Applicable to all the methods discussed so far and later
 - NAdam is popular