


`torch.compile()`

Eager Mode of Execution


- PyTorch by default runs models in **eager** mode
 - Dynamic execution
 - Prototyping: Easy to debug, highly flexible
 - Computation graph of code is created at runtime
 - Bad performance for production



Performance Gaps

Regime	What's happening	How to spot it
Compute-bound	GPU doing math	High SM utilisation 
Memory-bound	Shuffling tensors to DRAM	Low TFLOPS, high BW
Overhead-bound	Python dispatch & launches	CPU busy, GPU idle ← most models

Time breakdown for a typical transformer training step



Ref: <https://news.ycombinator.com/item?id=30697061>


Ref: https://horace.io/brrr_intro.html

```
import torch
from torch import profiler
from transformers import AutoModelForCausalLM, AutoTokenizer

model_name = "ibm-granite/granite-3.3-8b-instruct"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(model_name, dtype=torch.bfloat16).to("cuda")
inputs = tokenizer("What is the capital of France?", return_tensors="pt").to("cuda")

with torch.no_grad():
    with profiler.profile(
        activities=[profiler.ProfilerActivity.CPU, profiler.ProfilerActivity.CUDA],
        record_shapes=True,
        profile_memory=True,
        with_stack=True,
    ) as p:
        model(**inputs)

p.export_chrome_trace("trace-granite.json")
```




Area Selection Slices Pivot Table Slice Flamegraph Flow Events

Flat Tree

Name	Count	Wall Duration ↓	Self Duration	Wall Duration
	COUNT	SUM	SUM	AVG
ampere_bf16_s16816gemm_bf16_64x64_sliced1x2_ldg8_f2f_stages_64x6_tn	2	39.36µs	39.36µs	19.68µs
void at::native::elementwise_kernel<128, 4, at::native::gpu_kernel_impl_nocast<at::native::BinaryFu...	1	34.34µs	34.34µs	34.34µs
	1	5.025µs	5.025µs	5.025µs

131ms: GPU idle — CPU dispatching one Linear layer




Area Selection | Slices | Pivot Table | Slice Flamegraph | Flow Events

Export

Flat Tree

Name	Count	Wall Duration ↓		Self Duration		Wall Duration	
		COUNT	SUM	SUM	AVG	SUM	AVG
void at::native::reduce_kernel<512, 1, at::native::ReduceOp<float, at::native::MeanOps<float, float, ...	7	1	26.37µs	26.37µs	3.767µs	7.425µs	7.425µs
void at::native::elementwise_kernel<128, 4, at::native::gpu_kernel_impl_nocast<at::native::BinaryFu...	1	1	5.025µs	5.025µs	5.025µs	5.025µs	5.025µs
void at::native::elementwise_kernel<128, 2, at::native::gpu_kernel_impl_nocast<at::native::BinaryFu...	1	1	4.064µs	4.064µs	4.064µs	4.064µs	4.064µs
void at::native::vectorized_elementwise_kernel<4, at::native::bfloat16_copy_kernel_cuda(at::TensorI...	1	1	2.560µs	2.560µs	2.560µs	2.560µs	2.560µs
void at::native::vectorized_elementwise_kernel<4, at::native::(anonymous namespace)::pow_tensor...	1	1	2.560µs	2.560µs	2.560µs	2.560µs	2.560µs
void at::native::vectorized_elementwise_kernel<4, at::native::rsqrt_kernel_cuda(at::TensorIteratorBa...	1	1	2.464µs	2.464µs	2.464µs	2.464µs	2.464µs
void at::native::vectorized_elementwise_kernel<4, at::native::CUDAFunctionOnSelf_add<float>, std::a...	1	1	2.272µs	2.272µs	2.272µs	2.272µs	2.272µs

7 kernels, ~2-7µs each — one DRAM round-trip per op, no fusion



Slice ampere_bf16_s16816gemm_bf16_64x64_sliced1x2_ldg8_f2f_stages_64x6_tn

Contextual Options

Details

Name [ampere_bf16_s16816gemm_bf16_64x64_sliced1x2_ldg8_f2f_stages_64x6_tn](#)

Category kernel

Start time 00:00:00.433 520 046

Duration 73µs 216ns

Thread [stream 7 \[7\]](#)

Process [python3 \[0\]](#)

SQL ID [slice\[2201\]](#)

Preceding Flows

Slice	Delay	Thread
cudaLaunchKernel	1µs 42ns	thread 62238 (python3) 62238 (python3 62238)

Arguments

args

External id 233

queued 0

device 0

context 1


stream 7

correlation 4191

73µs GEMM — part of the 2.2% of time the GPU is actually computing

Graph Mode of Execution


- Graph mode enables the application of optimizations such as scheduling and operator fusion for faster execution/lower overheads



Enter torch.compile

- Just-in-Time compiler introduced in PyTorch 2.0
- No framework rewrite required, works with existing PyTorch code

```
@torch.compile
def foo(x: torch.Tensor, y: torch.Tensor):
    z = x * y
    w = z + x
    return w
```



Our Results — MobileNet v2

Ran the following:

- Model: MobileNet v2
- Dataset: ImageNet with 1000 sample images

Without Compile


```
Batch Size: 32  
Average Latency: 26.55 ± 0.28 ms  
Min Latency: 26.34 ms  
Max Latency: 27.74 ms  
Throughput: 1205.34 samples/sec
```

With Compile

```
Batch Size: 32  
Average Latency: 17.86 ± 0.17 ms  
Min Latency: 17.68 ms  
Max Latency: 18.33 ms  
Throughput: 1791.60 samples/sec
```


32% reduction in avg latency and 48% increase in throughput

What happens under the hood



Component I - Torch Dynamo


- Python level JIT compiler
 - Designed to allow graph compilation in PyTorch programs
 - Emits Fx Graphs containing sequences of PyTorch operations



Features of Torch Dynamo - Graph Breaks

- Unlike other tracers, torch dynamo does not crash when it encounters unknown code.
- It simply falls back to the CPython Interpreter to run the problematic code.

```
@torch.compile(backend=my_compiler)
def fn(x, y):
    a = torch.cos(x)
    b = torch.sin(y)
    return a + b
```



opcode	name	target	args	kwargs
placeholder	l_x_	L_x_	()	{}
placeholder	l_y_	L_y_	()	{}
call_function	a	<built-in method cos of type object at 0x7b565621ff00>	(l_x_,)	{}
call_function	b	<built-in method sin of type object at 0x7b565621ff00>	(l_y_,)	{}
call_function	add	<built-in function add>	(a, b)	{}
output	output	output	((add,),)	{}

Tabular view of FX graph

Features of Torch Dynamo - Graph Breaks

```
@torch.compile(backend=my_compiler)
def fn(x, y):
    a = torch.cos(x)
    b = torch.sin(y)
    print("Hello")
    return a + b
```

```
explanation = dynamo.explain(fn)(torch.randn(10), torch.randn(10))
print(explanation)
```


```
Hello
Graph Count: 2
Graph Break Count: 1
Op Count: 3
Break Reasons:
  Break Reason 1:
    Reason: builtin: print [<class 'torch._dynamo.variables.constant.ConstantVariable'>] False
  User Stack:
    <FrameSummary file /tmp/ipython-input-10-3243079242.py, line 12 in fn>
```

opcode	name	target	args	kwargs
placeholder	l_x_	L_x_	()	{}
placeholder	l_y_	L_y_	()	{}
call_function	a	<built-in method cos of type object at 0x7b565621ff00>	(l_x_,)	{}
call_function	b	<built-in method sin of type object at 0x7b565621ff00>	(l_y_,)	{}
output	output	output	((a, b),)	{}
Hello				
opcode	name	target	args	kwargs
placeholder	l_a_	L_a_	()	{}
placeholder	l_b_	L_b_	()	{}
call_function	add	<built-in function add>	(l_a_, l_b_)	{}
output	output	output	((add,),)	{}

Features of Torch Dynamo - Guards

Guards specify conditions (e.g., devices, data types, shapes, values) under which functions can be optimized

- If guard conditions satisfied, Dynamo replaces functions by their optimized versions in Python bytecode
- Otherwise, recompilation takes place under the new guard.



Features of Torch Dynamo - Guards

```
import torch

@torch.compile
def fn(x, n):
    y = x ** 2
    if n >= 0:
        return (n + 1) * y
    else:
        return y / n


x = torch.randn(200)
fn(x, 2)
fn(x, 3)
fn(x, -2)
```

```
uts.py
[guards] GUARDS:
[guards]
[guards] TREE_GUARD_MANAGER:
[guards] +- RootGuardManager
[guards] +- DEFAULT_DEVICE: utils_device.CURRENT_DEVICE == None # _dynamo/output_graph.py:493 in init_ambient_guards
[guards] +- GLOBAL_STATE: __check_global_state()
[guards] +- TORCH_FUNCTION_MODE_STACK: __check_torch_function_mode_stack()
[guards] +- GuardManager: source=L['n'], accessed_by=DictGetItemGuardAccessor('n')
[guards] | +- EQUALS_MATCH: L['n'] == 2 # if n >= 0: # sers/aishwariya/tuts.py:6 in fn
[guards] | +- GuardManager: source=L['x'], accessed_by=DictGetItemGuardAccessor('x')
[guards] | +- TENSOR_MATCH: check_tensor(L['x'], Tensor, DispatchKeySet(CPU, BackendSelect, ADInplaceOrView, AutogradCPU), torch.float32, device=None, re
uts.py:5 in fn
[guards] | | +- NO_HASATTR: hasattr(L['x'], '_dynamo_dynamic_indices') == False # y = x ** 2 # sers/aishwariya/tuts.py:5 in fn
[guards] Guard eval latency = 0.26 us
[recompiles] Recompiling function fn in /Users/aishwariya/tuts.py:3
[recompiles] triggered by the following guard failure(s):
[recompiles] - 0/0: L['n'] == 2 # if n >= 0: # sers/aishwariya/tuts.py:6 in fn
[guards] GUARDS:
[guards] TREE_GUARD_MANAGER:
[guards] +- RootGuardManager
[guards] +- DEFAULT_DEVICE: utils_device.CURRENT_DEVICE == None # _dynamo/output_graph.py:493 in init_ambient_guards
[guards] +- GLOBAL_STATE: __check_global_state()
[guards] +- TORCH_FUNCTION_MODE_STACK: __check_torch_function_mode_stack()
[guards] +- GuardManager: source=L['n'], accessed_by=DictGetItemGuardAccessor('n')
[guards] | +- TYPE_MATCH: __check_type_id(L['n'], 4361881088) # if n >= 0: # sers/aishwariya/tuts.py:6 in fn
[guards] | +- GuardManager: source=L['x'], accessed_by=DictGetItemGuardAccessor('x')
[guards] | +- TENSOR_MATCH: check_tensor(L['x'], Tensor, DispatchKeySet(CPU, BackendSelect, ADInplaceOrView, AutogradCPU), torch.float32, device=None, e
uts.py:5 in fn
[guards] | | +- NO_HASATTR: hasattr(L['x'], '_dynamo_dynamic_indices') == False # y = x ** 2 # sers/aishwariya/tuts.py:5 in fn
[guards] +- LAMBDA_GUARD: 0 <= L['n'] # if n >= 0: # sers/aishwariya/tuts.py:6 in fn (_dynamo/variables/tensor.py:1201 in evaluate_expr)
[guards] Guard eval latency = 0.28 us
[recompiles] Recompiling function fn in /Users/aishwariya/tuts.py:3
[recompiles] triggered by the following guard failure(s):
[recompiles] - 0/1: 0 <= L['n'] # if n >= 0: # sers/aishwariya/tuts.py:6 in fn (_dynamo/variables/tensor.py:1201 in evaluate_expr)
[recompiles] - 0/0: L['n'] == 2 # if n >= 0: # sers/aishwariya/tuts.py:6 in fn
[guards] GUARDS:
[guards] TREE_GUARD_MANAGER:
[guards] +- RootGuardManager
[guards] +- DEFAULT_DEVICE: utils_device.CURRENT_DEVICE == None # _dynamo/output_graph.py:493 in init_ambient_guards
[guards] +- GLOBAL_STATE: __check_global_state()
[guards] +- TORCH_FUNCTION_MODE_STACK: __check_torch_function_mode_stack()
[guards] +- GuardManager: source=L['n'], accessed_by=DictGetItemGuardAccessor('n')
[guards] | +- TYPE_MATCH: __check_type_id(L['n'], 4361881088) # if n >= 0: # sers/aishwariya/tuts.py:6 in fn
[guards] | +- GuardManager: source=L['x'], accessed_by=DictGetItemGuardAccessor('x')
[guards] | +- TENSOR_MATCH: check_tensor(L['x'], Tensor, DispatchKeySet(CPU, BackendSelect, ADInplaceOrView, AutogradCPU), torch.float32, device=None, re
uts.py:5 in fn
[guards] | | +- NO_HASATTR: hasattr(L['x'], '_dynamo_dynamic_indices') == False # y = x ** 2 # sers/aishwariya/tuts.py:5 in fn
[guards] +- LAMBDA_GUARD: L['n'] <= -1 # if n >= 0: # sers/aishwariya/tuts.py:6 in fn (_dynamo/variables/tensor.py:1201 in evaluate_expr)
[guards] Guard eval latency = 0.28 us
```

Recompilation
upon guard
failure

Component II - AOT Autograd

- Captures both forward and backward computations as FX graphs ahead of time
- It decides what values to save in order to run the forward and backward passes with minimal memory overheads and memory access latency



Component III - Inductor

- Default compiler backend for TorchDynamo
- Translates PyTorch programs into Triton for GPU and C++/OpenMP for CPU
- Performs different optimizations such as:

```
tmp1 = x + y  
tmp2 = tmp1 * z  
out = relu(tmp2)
```

Can be reduced to
out = relu((x + y) * z)

Operator fusion

```
y = sin(x)  
z = sin(x)  
return z
```

Can be reduced to
z = sin(x)
return z

Dead code elimination

```
scale = 2.0 * 4.0  
y = x * scale
```

Can be reduced to
y = x * 8.0

Constant folding


```
tmp = x + 1.0    y = tmp *  
2.0
```

Can be reduced to
y = x + 1.0 y *= 2.0

Buffer Reuse


Example of operator fusion

```
def foo(x: torch.Tensor, y: torch.Tensor):  
    z = x * y  
    w = z + x  
    return w
```



Example of operator fusion

```
@torch.compile
def foo(x: torch.Tensor, y: torch.Tensor):
    z = x * y
    w = z + x
    return w
```



Fused Kernel in Triton

```
triton_poi_fused_add_mul_0 = async_compile.triton('triton_', '''
...
@triton.jit
def triton_(in_ptr0, in_ptr1, out_ptr0, xnumel, XBLOCK : tl.constexpr):
    xnumel = 25
    xoffset = tl.program_id(0) * XBLOCK
    xindex = xoffset + tl.arange(0, XBLOCK)[: ]
    xmask = xindex < xnumel
    x0 = xindex
    tmp0 = tl.load(in_ptr0 + (x0), xmask)
    tmp1 = tl.load(in_ptr1 + (x0), xmask)
    tmp2 = tmp0 * tmp1
    tmp3 = tmp2 + tmp0
    tl.store(out_ptr0 + (x0), tmp3, xmask)
...

stream0 = get_raw_stream(0)
triton_poi_fused_add_mul_0.run(arg0_1, arg1_1, buf0, 25, grid=grid(25), stream=stream0)
del arg0_1
```

TODO: Try to generate the torch profiler trace for granite in compile mode and try to find out the differences.

Speed Up Numbers – Broader Benchmarks

TorchBench CPU (green=0 breaks · yellow=few · red=many)

Model	Speedup	Graph breaks
Background_Matting	1.51×	0
llama	1.36×	0
basic_gnn_edgecn	1.28×	0
BERT_pytorch	1.25×	0
basic_gnn_gcn	1.11×	6
alexnet	1.05×	0
basic_gnn_gin	1.04×	0
detectron2_fcoss_r_50_fpn	1.06×	20
detectron2_fasterrcnn_r_50_dc5	1.05×	32
detectron2_fasterrcnn_r_50_fpn	1.03×	34
detectron2_fasterrcnn_r_101_dc5	1.03×	32
detectron2_fasterrcnn_r_101_fpn	1.02×	34
detectron2_maskrcnn_r_101_c4	1.02×	44
detectron2_fasterrcnn_r_101_c4	0.97×	32
detectron2_fasterrcnn_r_50_c4	0.96×	32

Torch Compile vs Eager on Nvidia A100 GPUs

- ▶ TIMM models - 38% faster
- ▶ TorchBench models - 76% faster
- ▶ Huggingface models - 52% faster

Ref: Published on <https://pytorch.org/get-started/pytorch-2-x/>

Note: Detectron2 (32–44 breaks) still gets speedups.

Utilities - Modes, Configs

`torch.compile()` Arguments:

- `modes:`
 - `default`: efficient compilation with no extra memory and low compilation time
 - `reduce-overhead`: maximum framework overhead reduction, with little extra memory
 - `max-autotune`: fastest possible code, but long compilation time
- `fullgraph=True` → to create single graph with no graph break
- `dynamic=True` → all integers and shapes marked as dynamic

Other utilities:

- `torch._dynamo.utils.compile_times()` → to display time spent in each compilation phase
- `torch._dynamo.config.recompile_limit` → to limit the number of recompilations
- `torch._dynamo.explain()` → to display number of graph breaks and their causes
- `torch._dynamo.reset()` → to reset all artifacts created by dynamo
- `export TORCH_LOGS="graph_code,graph_sizes,dynamo,guards,output_code,graph_breaks"`

Deep Dive

1. https://docs.pytorch.org/docs/2.12/user_guide/torch_compiler/torch.compiler.html
2. <https://docs.google.com/document/d/1y5CRfMLdwEoF1nTk9q8qEu1mgMUuUtvhklPKJ2emLU8/edit?tab=t.0#heading=h.ivdr7fmrbeab>
3. <https://colab.research.google.com/drive/1Zh-Uo3TcTH8yYJF-LLo5rjlHVMtqvMdf?usp=sharing#scrollTo=ObOktQOeko5h>
4. https://docs.pytorch.org/docs/stable/torch.compiler_dynamo_deepdive.html
5. <https://docs.pytorch.org/assets/pytorch2-2.pdf>
6. <https://www.youtube.com/watch?v=rn-kJQ-7JmQ>
7. <https://www.youtube.com/watch?v=ppWKVg-VxmQ>
8. <https://www.youtube.com/watch?v=GmhnYe9QQoM>
9. <https://www.youtube.com/watch?v=5FNHwPlyHr8>
10. https://docs.google.com/document/d/1GgvOe7C8_NVOMLOCwDaYV1mXXyHMXy7ExoewHqooxrs/edit?tab=t.0#heading=h.fh8z zonyw8ng
11. <https://medium.com/data-science/how-pytorch-2-0-accelerates-deep-learning-with-operator-fusion-and-cpu-gpu-code-generation-35132a85bd26>