


Lecture 12: Off-Policy Evaluation in Bandits

ID 6002W: Online and Reinforcement Learning

Web-enabled M.Tech. program, IIT Madras

**How Do We Evaluate The Performance
Of A Bandit Algorithm In Practice?**

Regret Does Not Work In Practice



- In theory, and in simulations, we can calculate regret
- In practice, the best arm is unknown! $\Rightarrow$ regret impossible to calculate
- But we can measure cumulative reward for a bandit algorithm running in a real system

Practical Scenario

As a lead ML engineer

- You are running a bandit algorithm in production
 - e.g., recommending restaurants in Swiggy/Zomato

Practical Scenario

As a lead ML engineer

- You are running a bandit algorithm in production
 - e.g., recommending restaurants in Swiggy/Zomato
- Your team proposes several variants:
 - New reward model (lecture 11)
 - New way of designing context vectors (lecture 10)
 - New exploration-exploitation tradeoff constant (lecture 9)
- Which should you deploy?


A/B Testing

The Gold Standard

What Is an A/B Test?

- Split incoming traffic between two policies
 - Group A: current production policy
 - Group B: new candidate policy
- Run both for some duration

Recall:
adaptive stopping time $\leftrightarrow$
best-arm identification problem!




A/B Testing

The Gold Standard




What Is an A/B Test?

- Split incoming traffic between two policies
 - Group A: current production policy
 - Group B: new candidate policy
- Run both for some duration
- Observe rewards in each group, and compare them
 - decide whether to deploy new policy



A/B Testing Is Expensive


- A/B tests allocate real traffic to a candidate policy
- Highly risky and costly if candidate is worse!
 - May take many weeks
 - If new policy is very bad, users may never return!



A/B Testing Is Expensive

- A/B tests allocate real traffic to a candidate policy
- Highly risky and costly if candidate is worse!
 - May take many weeks
 - If new policy is very bad, users may never return!
- Cannot test every candidate online
 - Need an offline filter to triage candidate policies

How do we do that?



Off-Policy Evaluation Solves This Problem

Data For Off-Policy Evaluation

Production Logs

- A log records what the production policy did at every time step

t	external context $C(t)$	available articles $\mathcal{X}_t$	played $A(t)$	reward $y(t)$
1	user A asks: sports	{17, 3, 41, ...}	17	1
2	user B asks: tech	{12, 3, 28, ...}	3	0
3	user C asks: politics	{9, 14, 21, ...}	9	1
...	...	...	...	...

- Σ Each article $i \in \mathcal{X}_t$ has features $\mathbf{x}_i(t) \in \mathbb{R}^d$
- $\{\}$ Log of T impressions: $\{ (C(t), \mathcal{X}_t, A(t), \mathbf{y}(t)) \}_{t=1}^T$

Off-Policy Evaluation

- Suppose have logs from a policy that ran in production

Can we estimate, from these logs alone,
what a different policy would have done?

Off-Policy Evaluation

- Suppose have logs from a policy that ran in production

Can we estimate, from these logs alone,
what a different policy would have done?

This is called off-policy evaluation (OPE)

- Correct logging is super-critical!

Equally important in
reinforcement learning!

But What Is *A* Policy?

Defining A Policy

- A policy π is a rule for picking an arm from the action set $\mathcal{X} = \{x_1, x_2, \dots, x_K\}$
- $\pi(i \mid \mathcal{X})$: probability of picking arm i given $\mathcal{X}$

Defining A Policy

- A policy π is a rule for picking an arm from the action set $\mathcal{X} = \{x_1, x_2, \dots, x_K\}$
- $\pi(i \mid \mathcal{X})$: probability of picking arm i given $\mathcal{X}$
- A policy is fully specified by these probabilities
 - Uniform random: $\pi(i \mid \mathcal{X}) = 1/K$
 - Greedy on $\hat{\theta}$: $\pi(i \mid \mathcal{X}) = 1$ if $i = \arg \max \langle \hat{\theta}, x_i \rangle$
 - ε -greedy: $\pi(i \mid \mathcal{X}) = 1 - \varepsilon$ on $\arg \max \langle \hat{\theta}, x_i \rangle$, ε/K elsewhere

Defining A Policy

- A policy π is a rule for picking an arm from the action set $\mathcal{X} = \{x_1, x_2, \dots, x_K\}$
- $\pi(i \mid \mathcal{X})$: probability of picking arm i given $\mathcal{X}$
- A policy is fully specified by these probabilities
 - Deterministic policy: $\pi(i \mid \mathcal{X}) \in \{0, 1\}$
 - Randomised policy: $\pi(i \mid \mathcal{X})$ is a non-trivial distribution

Algorithm v/s Policy

- A learning algorithm induces a changing sequence of policies
 - Its internal state $(\hat{\theta}_t, V_t)$ updates as data comes in
 - $\pi(i \mid \mathcal{X})$ changes after each update

Algorithm v/s Policy

- A learning algorithm induces a changing sequence of policies
 - Its internal state $(\hat{\theta}_t, V_t)$ updates as data comes in
 - $\pi(i \mid \mathcal{X})$ changes after each update
- A frozen snapshot of LinUCB — fixed $(\hat{\theta}, V, \beta)$ — is a policy:
 - $\pi(i \mid \mathcal{X}) = 1$ if $i = \arg \max_j \langle \hat{\theta}, x_j \rangle + \beta \|x_j\|_{V^{-1}}$, else 0
 - Deterministic, assuming fixed tie-breaking

Algorithm v/s Policy

- A learning algorithm induces a changing sequence of policies
 - Its internal state $(\hat{\theta}_t, V_t)$ updates as data comes in
 - $\pi(i \mid \mathcal{X})$ changes after each update
- A frozen snapshot of LinTS — fixed posterior over θ — is a policy:
 - Sample $\tilde{\theta} \sim$ posterior, play $\arg \max \langle \tilde{\theta}, x_i \rangle$
 - Randomised, because $\tilde{\theta}$ is sampled fresh each time

Off-Policy Evaluation

Setting and Notation

- In off-policy evaluation, we have two policies
 - Logging policy: $\pi_0(i \mid \mathcal{X})$ — produced the data Called logger for short
 - Target policy: $\pi(i \mid \mathcal{X})$ — a new policy we want to evaluate
 - For simplicity, we will assume they are fixed policies, not learning algorithms

Off-Policy Evaluation

Setting and Notation

- In off-policy evaluation, we have two policies
 - Logging policy: $\pi_0(i \mid \mathcal{X})$ — produced the data Called logger for short
 - Target policy: $\pi(i \mid \mathcal{X})$ — a new policy we want to evaluate
 - For simplicity, we will assume they are fixed policies, not learning algorithms
- We care about value of a policy: expected reward in one time step if π is used

$$V(\pi) = \mathbb{E}_{\mathcal{X}}[\sum \pi(i \mid \mathcal{X}) \langle \theta, x_i \rangle]$$


- Goal: estimate $V(\pi)$ using logs collected under π_0

The Counterfactual Question

What we observe versus what we want

We observe what the logging policy did.

We want to know what a different policy would have done.



Same impression, two policies, only one observation.

The Replay Method

When The Logger Is Uniform

- Suppose $\pi_0(i | \mathcal{X}) = 1/K$ for every arm i (uniform exploration)
- Suppose target π is deterministic: in each context, π picks one arm

When The Logger Is Uniform

- Suppose $\pi_0(i \mid \mathcal{X}) = 1/K$ for every arm i (uniform exploration)
 - Suppose target π is deterministic: in each context, π picks one arm
- ✱ Idea: go step by step along the log
- ➡ Keep the rows where the logger happened to pick what π would pick
 - ➡ Discard the rest

When The Logger Is Uniform

- Suppose $\pi_0(i | \mathcal{X}) = 1/K$ for every arm i (uniform exploration)
 - Suppose target π is deterministic: in each context, π picks one arm
- ✱ Idea: go step by step along the log
- ➡ Keep the rows where the logger happened to pick what π would pick
 - ➡ Discard the rest
 - Introduced by Li et al. (2010) for evaluating news recommendation on Yahoo!
 - a small slice of traffic served by uniform random policy, kept aside specifically for evaluation

Replay Method Mechanics

Simplest Off-Policy Method



Replay keeps only the rows where **logger** and **target** agree.

Analysing Replay Method

Why does this method work?

- The logger chooses each arm uniformly at random
 - The event “logger agrees with π ” occurs with probability $1/K$, independent of reward
 - Kept rows are a random sample of the times where π was effectively played
 - Averaging their rewards gives an estimate of $V(\pi)$

Analysing Replay Method

Why does this method work?

- The logger chooses each arm uniformly at random
 - The event “logger agrees with π ” occurs with probability $1/K$, independent of reward
 - Kept rows are a random sample of the times where π was effectively played
 - Averaging their rewards gives an estimate of $V(\pi)$
- Cons:
 - Only $1/K$ of the log is used $\rightarrow$ wasteful for large K
 - Requires uniform logging $\rightarrow$ does not apply to production policies

We need a method that works for any known logging policy π_0

The Inverse-Propensity Scoring (IPS) Method

What Must Be Logged

Per impression/interaction/round t , record:

What Must Be Logged

Per impression/interaction/round t , record:

- Action set: $\mathcal{X}(t) = \{x_1(t), \dots, x_K(t)\}$
- Played arm: $A(t) \in \{1, \dots, K\}$
- Reward: $y(t)$

What Must Be Logged

Per impression/interaction/round t , record:

- Action set: $\mathcal{X}(t) = \{x_1(t), \dots, x_K(t)\}$
- Played arm: $A(t) \in \{1, \dots, K\}$
- Reward: $y(t)$
- Propensity: $p(t) = \pi_0(A(t) \mid \mathcal{X}(t))$

Note that propensity cannot be recovered from other data!

The IPS Estimator

Estimator

$$\hat{V}_{\text{IPS}}(\pi) = \frac{1}{T} \sum_{t=1}^T \frac{\pi(A_t \mid \mathcal{X}_t)}{p_t} y_t, \quad p_t = \pi_0(A_t \mid \mathcal{X}_t).$$

The IPS Estimator

Estimator

$$\hat{V}_{\text{IPS}}(\pi) = \frac{1}{T} \sum_{t=1}^T \frac{\pi(A_t | \mathcal{X}_t)}{p_t} y_t, \quad p_t = \pi_0(A_t | \mathcal{X}_t).$$

Why the ratio works: condition on one impression $\mathcal{X}$

$$\begin{aligned} \mathbb{E}_{A \sim \pi_0} \left[\frac{\pi(A | \mathcal{X})}{\pi_0(A | \mathcal{X})} \mu(\mathcal{X}, A) \right] &= \sum_{i \in \mathcal{X}} \pi_0(i | \mathcal{X}) \frac{\pi(i | \mathcal{X})}{\pi_0(i | \mathcal{X})} \mu(\mathcal{X}, i) \\ &= \sum_{i \in \mathcal{X}} \pi(i | \mathcal{X}) \mu(\mathcal{X}, i) \\ &= V(\pi | \mathcal{X}). \end{aligned}$$

The IPS Estimator

Estimator

$$\hat{V}_{\text{IPS}}(\pi) = \frac{1}{T} \sum_{t=1}^T \frac{\pi(A_t | \mathcal{X}_t)}{p_t} y_t, \quad p_t = \pi_0(A_t | \mathcal{X}_t).$$

Why the ratio works: condition on one impression $\mathcal{X}$

$$\begin{aligned} \mathbb{E}_{A \sim \pi_0} \left[\frac{\pi(A | \mathcal{X})}{\pi_0(A | \mathcal{X})} \mu(\mathcal{X}, A) \right] &= \sum_{i \in \mathcal{X}} \pi_0(i | \mathcal{X}) \frac{\pi(i | \mathcal{X})}{\pi_0(i | \mathcal{X})} \mu(\mathcal{X}, i) \\ &= \sum_{i \in \mathcal{X}} \pi(i | \mathcal{X}) \mu(\mathcal{X}, i) \\ &= V(\pi | \mathcal{X}). \end{aligned}$$

Takeaway. After averaging over impressions, $\mathbb{E}[\hat{V}_{\text{IPS}}(\pi)] = V(\pi)$, provided $\pi_0(i | \mathcal{X}) > 0$ whenever $\pi(i | \mathcal{X}) > 0$.

The Variance Cost

- IPS is unbiased — but variance can be enormous!
- If $p(t)$ is very small for some t , the weight $\pi(A(t) | x(t))/p(t)$ becomes very large
 - A single such sample can dominate the estimate!
 - Variance grows with square of the weight
- Practical cost of going from replay to IPS
 - We use all the data, but the price is variance
- Many heuristics to reduce variance exist (we won't cover)

Conclusion

- New bandit algorithms need to be safely evaluated before deployment!
- Off-policy evaluation: estimate value of a policy from logs of a different policy
- Requires:
 - a stochastic logger
 - logged propensities
- Estimator: Inverse Propensity Scoring - Need to watch out for high variance!
- OPE is used as a filter before A/B tests, not a replacement for them